



## Question 1

### Defining the problem statement

The Management team at Walmart Inc. wants to analyze the customer purchase behavior (specifically, purchase amount) against the customer's gender and the various other factors to help the business make better decisions. They want to understand if the spending habits differ between male and female customers: Do women spend more on Black Friday than men? (Assume 50 million customers are male and 50 million are female).

In [1018]:

```
##importing the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import norm
```

In [1019]:

```
##Loading the data
walmart = pd.read_csv("walmart.csv")
walmart.head()
```

Out[1019]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category	Purchas
0	1000001	P00069042	F	0-17	10	A	2	0	3	837
1	1000001	P00248942	F	0-17	10	A	2	0	1	1520
2	1000001	P00087842	F	0-17	10	A	2	0	12	142
3	1000001	P00085442	F	0-17	10	A	2	0	12	105
4	1000002	P00285442	M	55+	16	C	4+	0	8	796

### Analyzing the basic metrics

## Observations on shape of data

In [1020]:

```
walmart.shape ## 550068 rows and 10 columns
```

Out[1020]:

```
(550068, 10)
```

In [1021]:

```
walmart.ndim
```

Out[1021]:

```
2
```

In [1022]:

```
walmart.size
```

Out[1022]:

```
5500680
```

## data types of all the attributes

In [1023]:

```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   User_ID               550068 non-null  int64
 1   Product_ID           550068 non-null  object
 2   Gender               550068 non-null  object
 3   Age                  550068 non-null  object
 4   Occupation           550068 non-null  int64
 5   City_Category        550068 non-null  object
 6   Stay_In_Current_City_Years  550068 non-null  object
 7   Marital_Status       550068 non-null  int64
 8   Product_Category     550068 non-null  int64
 9   Purchase             550068 non-null  int64
dtypes: int64(5), object(5)
memory usage: 42.0+ MB
```

- User\_ID, Occupation, Marital\_Status, Product\_Category, Purchase are numerical
- Rest all are non numerical or categorical in the original data
- No **null** values

## conversion of categorical attributes to 'category'

In [1024]:

```
walmart["Marital_Status"] = walmart["Marital_Status"].astype('object')
```

In [1025]:

```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                            550068 non-null  object
2   Gender                                550068 non-null  object
3   Age                                    550068 non-null  object
4   Occupation                             550068 non-null  int64
5   City_Category                         550068 non-null  object
6   Stay_In_Current_City_Years            550068 non-null  object
7   Marital_Status                        550068 non-null  object
8   Product_Category                      550068 non-null  int64
9   Purchase                              550068 non-null  int64
dtypes: int64(4), object(6)
memory usage: 42.0+ MB
```

statistical summary

In [1026]:

```
walmart.describe() ##numerical attributes
```

Out[1026]:

	User_ID	Occupation	Product_Category	Purchase
count	5.500680e+05	550068.000000	550068.000000	550068.000000
mean	1.003029e+06	8.076707	5.404270	9263.968713
std	1.727592e+03	6.522660	3.936211	5023.065394
min	1.000001e+06	0.000000	1.000000	12.000000
25%	1.001516e+06	2.000000	1.000000	5823.000000
50%	1.003077e+06	7.000000	5.000000	8047.000000
75%	1.004478e+06	14.000000	8.000000	12054.000000
max	1.006040e+06	20.000000	20.000000	23961.000000

In [1027]:

```
walmart[["User_ID","Occupation","Product_Category","Purchase"]].mean()
```

Out[1027]:

```
User_ID          1.003029e+06
Occupation        8.076707e+00
Product_Category  5.404270e+00
Purchase          9.263969e+03
dtype: float64
```

In [1028]:

```
walmart[["User_ID","Occupation","Product_Category","Purchase"]].median()
```

Out[1028]:

```
User_ID          1003077.0
Occupation         7.0
Product_Category   5.0
Purchase          8047.0
dtype: float64
```

Occupation and Purchase have outliers present since mean and median have considerable difference

In [1029]:

```
walmart.mode()
```

Out[1029]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category	Purchase
0	1001680	P00265242	M	26-35	4	B	1	0	5	701

- Most **frequent buyer** is 1001680
- Most of the buyers are **male** and hail from **city category B**
- Most of the buyers have **occupation 4**
- Most of the buyers are **unmarried**
- **Product\_category 5** with product\_id P00265242 is most frequently bought

## Non-Graphical Analysis: Value counts and unique attributes

In [1030]:

```
walmart.nunique()
```

Out[1030]:

User_ID	5891
Product_ID	3631
Gender	2
Age	7
Occupation	21
City_Category	3
Stay_In_Current_City_Years	5
Marital_Status	2
Product_Category	20
Purchase	18105
dtype:	int64

In [1031]:

```
walmart["User_ID"].value_counts()
```

Out[1031]:

1001680	1026
1004277	979
1001941	898
1001181	862
1000889	823
...	
1002690	7
1002111	7
1005810	7
1004991	7
1000708	6
Name: User_ID, Length: 5891, dtype: int64	

1001680 is the most frequent buyer

In [1032]:

```
walmart["Product_ID"].value_counts()
```

Out[1032]:

```
P00265242    1880
P00025442     1615
P00110742     1612
P00112142     1562
P00057642     1470
...
P00314842         1
P00298842         1
P00231642         1
P00204442         1
P00066342         1
Name: Product_ID, Length: 3631, dtype: int64
```

- P00265242 is the most common product

In [1033]:

```
walmart["Product_ID"].nunique()
```

Out[1033]:

```
3631
```

- Only 3631 product\_id's are unique

In [1034]:

```
walmart["Gender"].value_counts()
```

Out[1034]:

```
M    414259
F    135809
Name: Gender, dtype: int64
```

- Most of the buyers are Males

In [1035]:

```
walmart["Age"].value_counts()
```

Out[1035]:

```
26-35    219587
36-45    110013
18-25     99660
46-50     45701
51-55     38501
55+       21504
0-17      15102
Name: Age, dtype: int64
```

- Most of the buyers are in Age group 26 to 35 and 36-45

In [1036]:

```
walmart["Occupation"].value_counts()
```

Out[1036]:

```
4      72308
0      69638
7      59133
1      47426
17     40043
20     33562
12     31179
14     27309
2      26588
16     25371
6      20355
3      17650
10     12930
5      12177
15     12165
11     11586
19      8461
13      7728
18      6622
9       6291
8       1546
Name: Occupation, dtype: int64
```

- Most of the buyers have occupation 4

In [1037]:

```
walmart["City_Category"].value_counts()
```

Out[1037]:

```
B      231173
C      171175
A      147720
Name: City_Category, dtype: int64
```

- Most of the buyers stay in city category B

In [1038]:

```
walmart["Stay_In_Current_City_Years"].value_counts()
```

Out[1038]:

```
1      193821
2      101838
3       95285
4+      84726
0       74398
Name: Stay_In_Current_City_Years, dtype: int64
```

- Most number of buyers have been staying for less than 2 years in the current city

In [1039]:

```
walmart["Marital_Status"].value_counts()
```

Out[1039]:

```
0      324731
1      225337
Name: Marital_Status, dtype: int64
```

- Mostly buyers are unmarried

In [1040]:

```
walmart["Product_Category"].value_counts()
```

Out[1040]:

```
5      150933
1      140378
8      113925
11     24287
2      23864
6      20466
3      20213
4      11753
16     9828
15     6290
13     5549
10     5125
12     3947
7       3721
18     3125
20     2550
19     1603
14     1523
17       578
9        410
```

Name: Product\_Category, dtype: int64

- Product category 5 is the most bought along with 1,8

In [1041]:

```
walmart["Purchase"].value_counts()
```

Out[1041]:

```
7011      191
7193      188
6855      187
6891      184
7012      183
...
23491      1
18345      1
3372       1
855        1
21489      1
```

Name: Purchase, Length: 18105, dtype: int64

In [1042]:

```
walmart["Purchase"].value_counts()[:50]
```

Out[1042]:

7011	191
7193	188
6855	187
6891	184
7012	183
6960	183
6879	182
7166	182
7027	182
6868	180
7165	180
6883	180
6858	179
7093	178
6931	178
7089	178
7185	178
6923	178
7114	177
7188	177
7085	176
6908	176
7060	176
7167	175
6973	175
6928	175
6949	175
7146	175
7159	175
6904	174
7010	174
7962	174
6952	174
7192	174
7034	174
6862	173
7047	173
7067	172
7108	172
7049	172
6930	172
7081	172
7028	172
6978	172
7110	171
6938	171
7024	171
7046	171
7026	171
7083	171

Name: Purchase, dtype: int64

- Most of the purchase are in the range 6k to 7k



Visual Analysis - Univariate & Bivariate

Univariate

In [1043]:

```
walmart.head()
```

Out[1043]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category	Purchas
0	1000001	P00069042	F	0-17	10	A	2	0	3	837
1	1000001	P00248942	F	0-17	10	A	2	0	1	1520
2	1000001	P00087842	F	0-17	10	A	2	0	12	142
3	1000001	P00085442	F	0-17	10	A	2	0	12	105
4	1000002	P00285442	M	55+	16	C	4+	0	8	796

In [1044]:


```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                        550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  object
7   Marital_Status                      550068 non-null  object
8   Product_Category                    550068 non-null  int64
9   Purchase                            550068 non-null  int64
dtypes: int64(4), object(6)
memory usage: 42.0+ MB
```

## Continuous Univariate

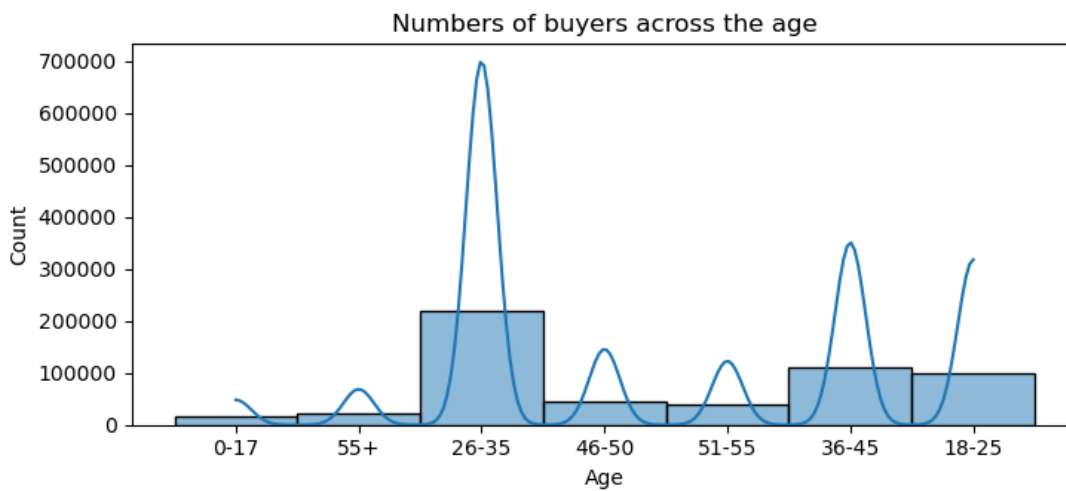
In [1045]:

```
plt.rcParams["figure.figsize"] = [7.50, 3.50]
plt.rcParams["figure.autolayout"] = True
sns.histplot(data=walmart["User_ID"],bins=35,kde=True)
plt.show()
```



In [1046]:

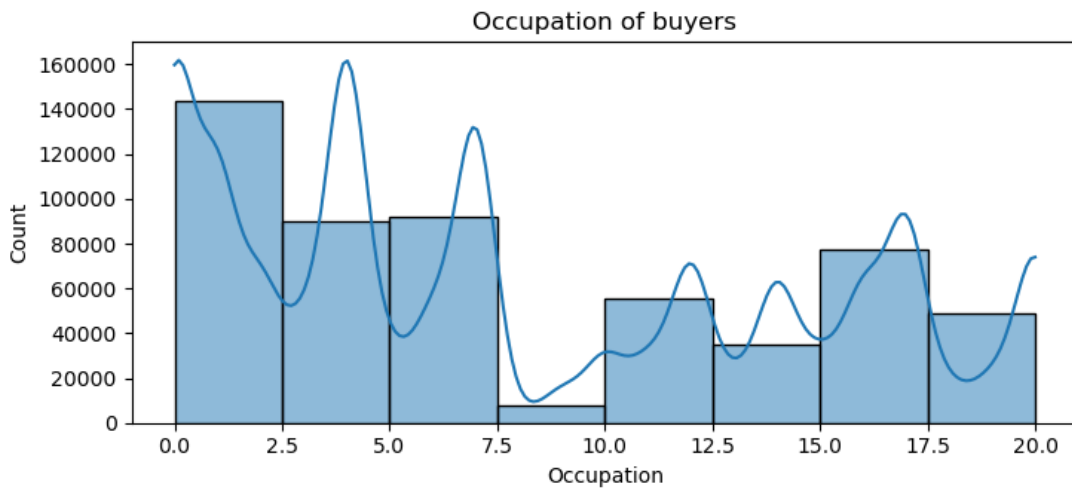
```
plt.rcParams["figure.figsize"] = [7.50, 3.50]
plt.rcParams["figure.autolayout"] = True
sns.histplot(data=walmart["Age"],bins=14,kde=True)
plt.title("Numbers of buyers across the age")
plt.show()
```



- Most of the users are between 26-35 & 36-45

In [1047]:

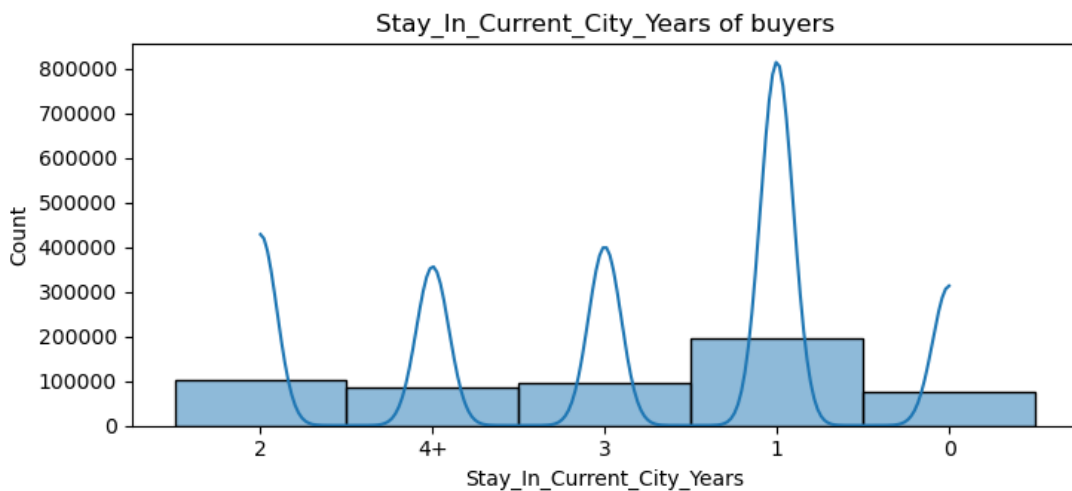
```
plt.rcParams["figure.figsize"] = [7.50, 3.50]
plt.rcParams["figure.autolayout"] = True
sns.histplot(data=walmart["Occupation"],binwidth=2.5,kde=True)
plt.title("Occupation of buyers")
plt.show()
```



- Most of the buyers belong to occupation 4

In [1048]:

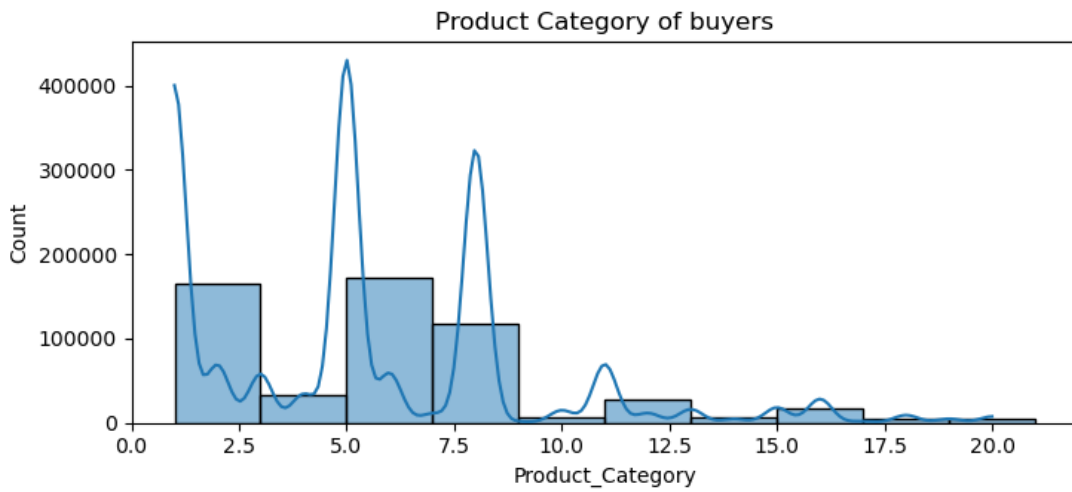
```
plt.rcParams["figure.figsize"] = [7.50, 3.50]
plt.rcParams["figure.autolayout"] = True
sns.histplot(data=walmart["Stay_In_Current_City_Years"],kde=True)
plt.title("Stay_In_Current_City_Years of buyers")
plt.show()
```



- Most of buyers have stayed only for 1 year in the current city
- Very less buyers have been staying in the city for less than a year

In [1049]:


```
plt.rcParams["figure.figsize"] = [7.50, 3.50]
plt.rcParams["figure.autolayout"] = True
sns.histplot(data=walmart["Product_Category"],binwidth=2,kde=True)
plt.title("Product Category of buyers")
plt.show()
```



- Product Category 5 is the most bought

In [1050]:

```
plt.rcParams["figure.figsize"] = [7.50, 3.50]
plt.rcParams["figure.autolayout"] = True
sns.histplot(data=walmart["Purchase"],binwidth=500,kde=True)
plt.title("Purchase of buyers")
plt.show()
```



- Most of the buyers spent within 5000 to 12000

## Categorical Univariate


In [1051]:

```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   User_ID             550068 non-null  int64  
 1   Product_ID          550068 non-null  object  
 2   Gender              550068 non-null  object  
 3   Age                 550068 non-null  object  
 4   Occupation          550068 non-null  int64  
 5   City_Category       550068 non-null  object  
 6   Stay_In_Current_City_Years  550068 non-null  object  
 7   Marital_Status      550068 non-null  object  
 8   Product_Category    550068 non-null  int64  
 9   Purchase            550068 non-null  int64  
dtypes: int64(4), object(6)
memory usage: 42.0+ MB
```

In [1052]:


```
sns.countplot(data=walmart,x="Product_Category",order = walmart['Product_Category'].value_counts().index)
plt.xticks(rotation=90,fontsize=12) ## to avoid overlapping labels
plt.show()
```



- Product Category 5,1,8 are the most bought

In [1053]:

```
sns.countplot(data=walmart,x="Gender",order = walmart['Gender'].value_counts().index)
plt.xticks(rotation=90,fontsize=12) ## to avoid overlapping labels
plt.show()
```



In [1054]:

```
walmart.groupby("Gender").count()["User_ID"]
```

Out[1054]:

```
Gender
F    135809
M    414259
Name: User_ID, dtype: int64
```

In [1055]:

```
(135809/(135809+414259))*100
```

Out[1055]:

```
24.689492935418894
```

In [1056]:

```
414259/(135809+414259)
```


Out[1056]:

```
0.7531050706458111
```

- Only 24 % of buyers are females, and 75% buyers are males

In [1057]:

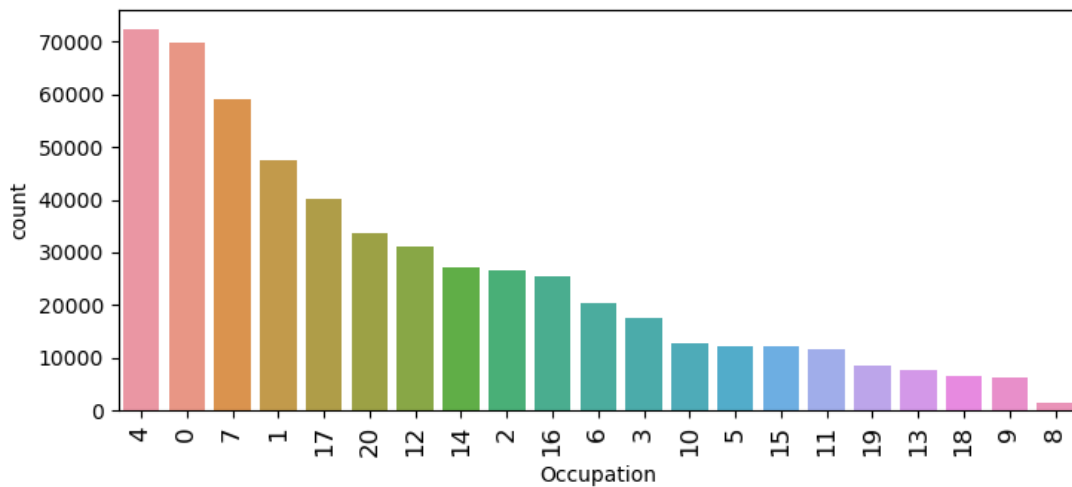
```
sns.countplot(data=walmart,x="Age",order = walmart['Age'].value_counts().index)
plt.xticks(rotation=90,fontsize=12) ## to avoid overlapping labels
plt.show()
```



- 26-25 age group has most number of buyers

In [1058]:

```
sns.countplot(data=walmart,x="Occupation",order = walmart['Occupation'].value_counts().index)
plt.xticks(rotation=90,fontsize=12) ## to avoid overlapping labels
plt.show()
```



- Most of the buyers have occupation 4 or occupation 0

In [1059]:

```
walmart.loc[walmart["Occupation"]==0]["Gender"].value_counts()
```

Out[1059]:

```
M    51526
F    18112
Name: Gender, dtype: int64
```

In [1060]:


```
walmart.loc[walmart["Occupation"]==4]["Gender"].value_counts()
```

Out[1060]:

```
M    54472
F    17836
Name: Gender, dtype: int64
```

In [1061]:


```
sns.countplot(data=walmart,x="Stay_In_Current_City_Years",order = walmart['Stay_In_Current_City_Years'].value_counts().index)
plt.xticks(rotation=90,fontsize=12) ## to avoid overlapping labels
plt.show()
```



- Most of the buyers are have been living in the current city for 1-2 years

In [1062]:

```
sns.countplot(data=walmart,x="City_Category",order = walmart['City_Category'].value_counts().index)
plt.xticks(rotation=90,fontsize=12) ## to avoid overlapping labels
plt.show()
```




- Most of the belong to city category B



In [1063]:

```
sns.countplot(data=walmart,x="Marital_Status",order = walmart['Marital_Status'].value_counts().index)
plt.xticks(rotation=90,fontsize=12) ## to avoid overlapping labels
plt.show()
```



- Most of the buyers are unmarried

In [1064]:

```
walmart.groupby("Marital_Status").count()["User_ID"]
```

Out[1064]:

```
Marital_Status
0      324731
1      225337
Name: User_ID, dtype: int64
```

In [1065]:

```
324731/(324731+225337)
```

Out[1065]:

```
0.5903470116421969
```

In [1066]:

```
225337/(324731+225337)
```

Out[1066]:

```
0.40965298835780306
```

- 59% of buyers are married 40% are unmarried

In [1067]:

```
pd.crosstab(index=walmart['Gender'], columns=[walmart['Marital_Status']],normalize=True,margins=True)
```


Out[1067]:

Marital_Status	0	1	All
Gender			
F	0.143293	0.103602	0.246895
M	0.447054	0.306051	0.753105
All	0.590347	0.409653	1.000000

### Top 10 product categories

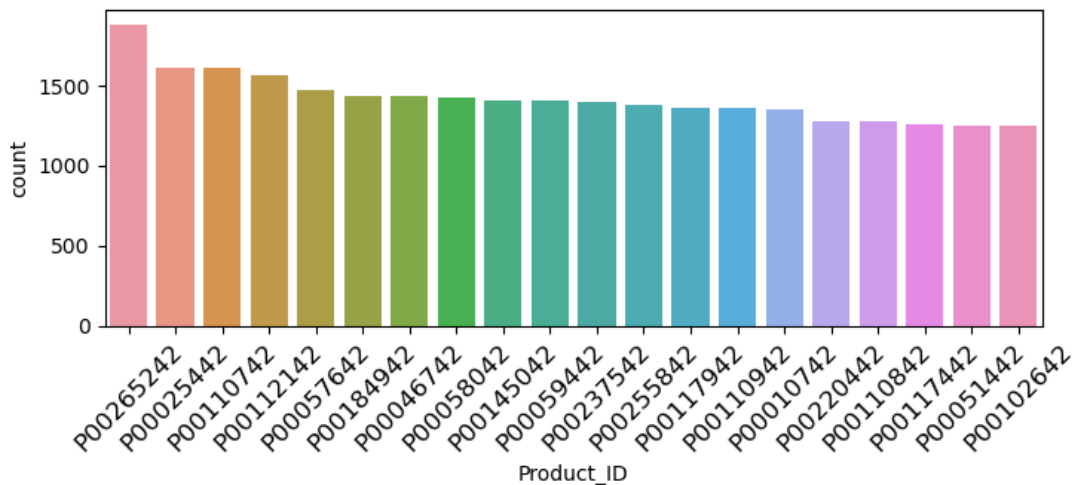
In [1068]:

```
walmart_top10_product_category = walmart["Product_Category"].value_counts()[:10]
top10_productcategory=walmart[walmart["Product_Category"].isin(walmart_top10_product_category.index)]
sns.countplot(data=top10_productcategory,x="Product_Category",order = top10_productcategory['Product_Category'].value_counts().index)
plt.xticks(rotation=45,fontsize=12) ## to avoid overlapping labels
plt.show()
```



In [1069]:

```
##Top 20 product_ID
walmart_top20_product_id = walmart["Product_ID"].value_counts()[:20]
top20_productid=walmart[walmart["Product_ID"].isin(walmart_top20_product_id.index)]
sns.countplot(data=top20_productid,x="Product_ID",order = top20_productid['Product_ID'].value_counts().index)
plt.xticks(rotation=45,fontsize=12) ## to avoid overlapping labels
plt.show()
```



- 'P00265242', 'P00025442', 'P00110742' are most sold product\_id's

## Bivariate

### Bivariate Analysis of two categorical variables

In [1070]:

```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   User_ID               550068 non-null  int64
 1   Product_ID           550068 non-null  object
 2   Gender               550068 non-null  object
 3   Age                  550068 non-null  object
 4   Occupation           550068 non-null  int64
 5   City_Category        550068 non-null  object
 6   Stay_In_Current_City_Years  550068 non-null  object
 7   Marital_Status        550068 non-null  object
 8   Product_Category      550068 non-null  int64
 9   Purchase              550068 non-null  int64
dtypes: int64(4), object(6)
memory usage: 42.0+ MB
```

### Top product\_category

In [1071]:

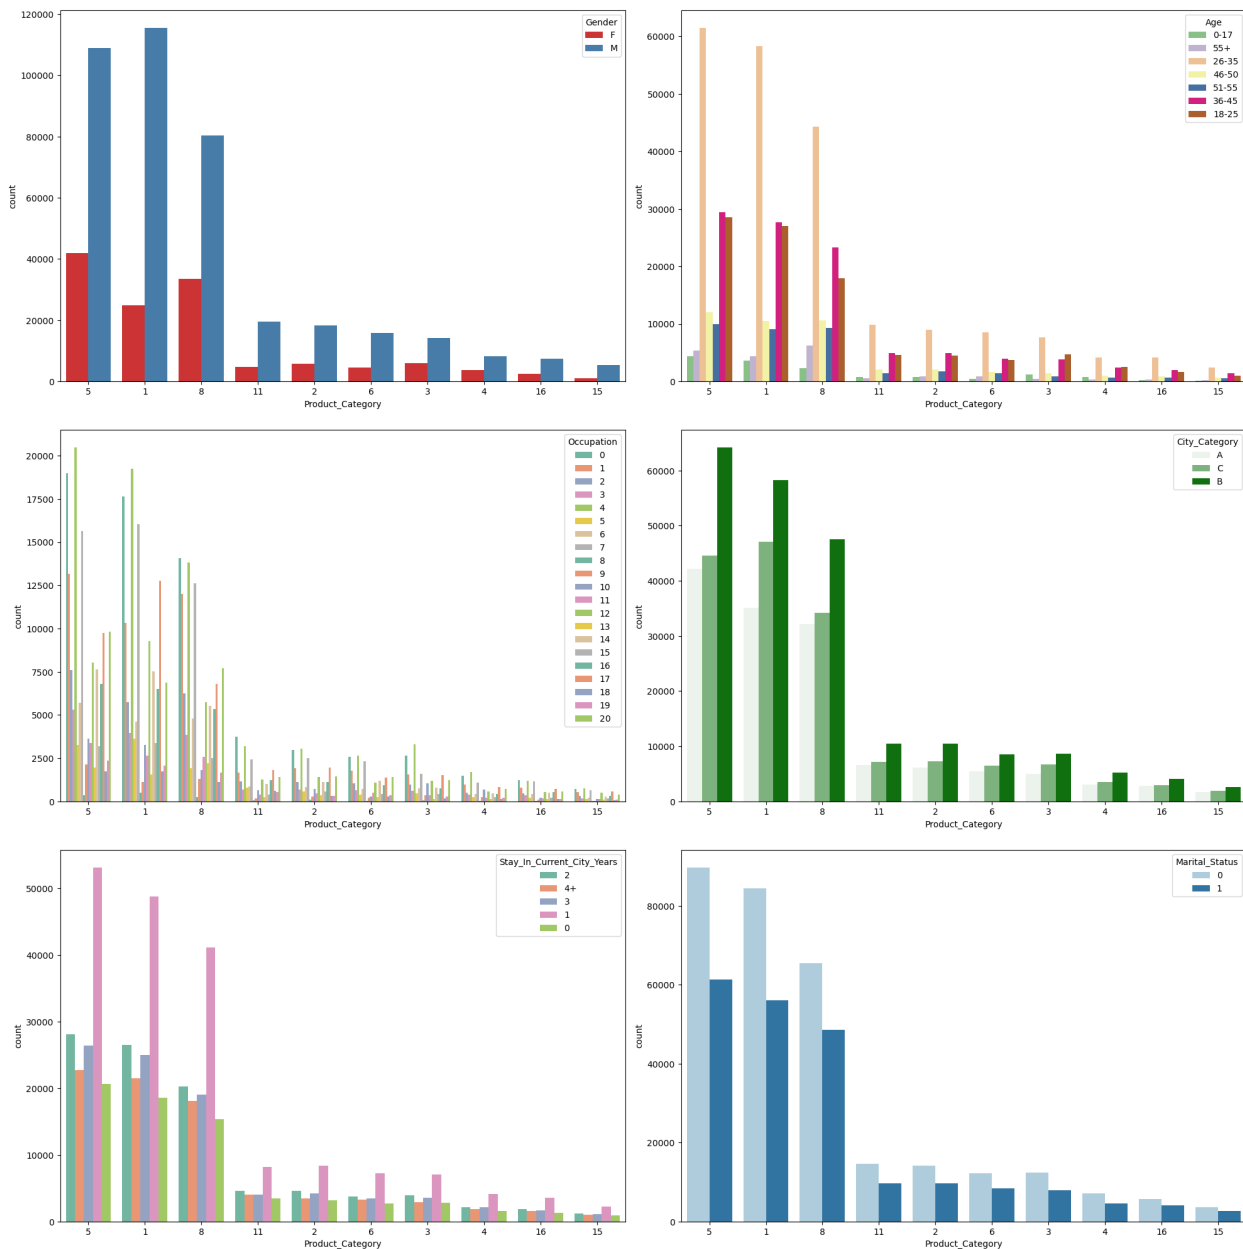
```
##Top 10 product_category
walmart_top10_product_category = walmart["Product_Category"].value_counts()[:10]
top10_productcategory=walmart[walmart["Product_Category"].isin(walmart_top10_product_category.index)]
```

In [1072]:

```

plt.figure(figsize=(20,20))
plt.subplot(3,2,1) ##1 shows the position
sns.countplot(data=top10_productcategory,x='Product_Category',hue='Gender',order=top10_productcategory['Product_Category'].value)
plt.subplot(3,2,2)
sns.countplot(data=top10_productcategory,x='Product_Category',hue='Age',order=top10_productcategory['Product_Category'].value)
plt.subplot(3,2,3)
sns.countplot(data=top10_productcategory,x='Product_Category',hue='Occupation',order=top10_productcategory['Product_Category'].value)
plt.subplot(3,2,4)
sns.countplot(data=top10_productcategory,x='Product_Category',hue='City_Category',order=top10_productcategory['Product_Category'].value)
plt.subplot(3,2,5)
sns.countplot(data=top10_productcategory,x='Product_Category',hue='Stay_In_Current_City_Years',order=top10_productcategory['Product_Category'].value)
plt.subplot(3,2,6)
sns.countplot(data=top10_productcategory,x='Product_Category',hue='Marital_Status',order=top10_productcategory['Product_Category'].value)
plt.show()


```



- Most of the buyers across the top 10 product category are males, female buyers are very less comparably
- Most of the buyers are between 26-35 years of age, very less buyers in age group 0 to 17 years of age
- Most of the buyers have occupation 4, occupation 8 buyers are very few
- Most of the buyers for these product category belong to city category B
- Most of the buyers for these product category have been the city for 0-1 years
- Most of the buyers are unmarried

In [1073]:

```
plt.figure(figsize=(20,10))
sns.countplot(data=top20_productid,x='Product_ID',hue='Product_Category',order=top20_productid['Product_ID'].value_counts())
plt.show()
```



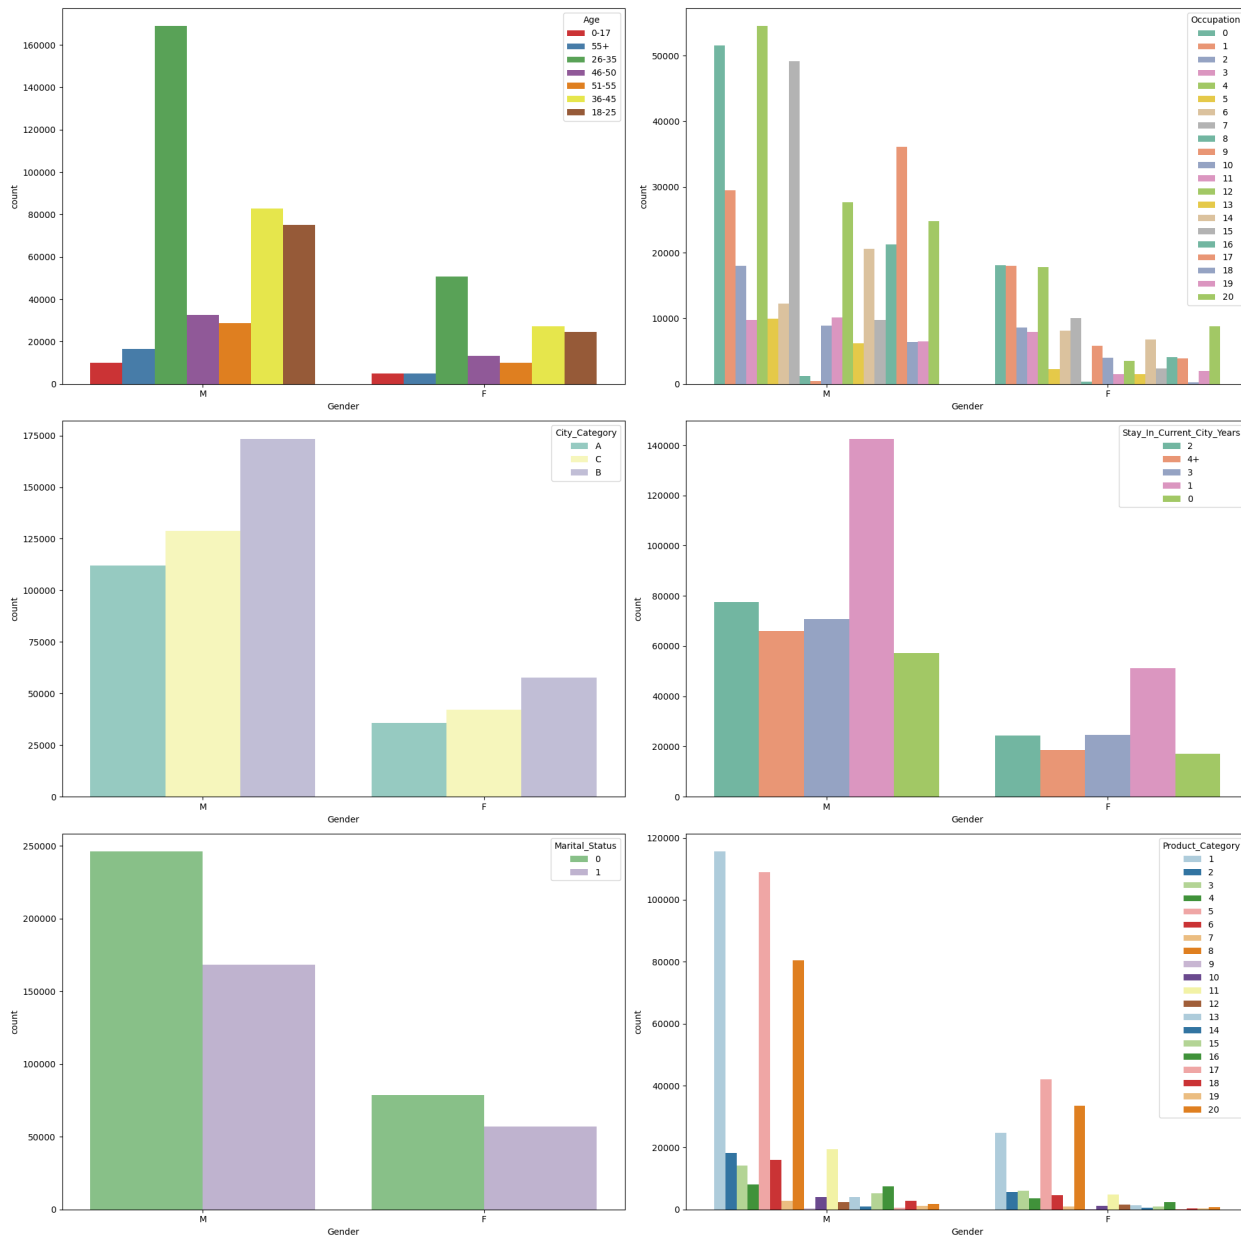
- Product category 5 is the most popular amongst top 5 productids
- Product category 1 is common for most of the product\_id and is also more frequently bought along with 5

In [1074]:

```

plt.figure(figsize=(20,20))
plt.subplot(3,2,1) ##1 shows the position
sns.countplot(data=walmart,x='Gender',hue='Age',order=walmart['Gender'].value_counts().index,palette="Set1")
plt.subplot(3,2,2)
sns.countplot(data=walmart,x='Gender',hue='Occupation',order=walmart['Gender'].value_counts().index,palette="Set2")
plt.subplot(3,2,3)
sns.countplot(data=walmart,x='Gender',hue='City_Category',order=walmart['Gender'].value_counts().index,palette="Set3")
plt.subplot(3,2,4)
sns.countplot(data=walmart,x='Gender',hue='Stay_In_Current_City_Years',order=walmart['Gender'].value_counts().index,palette="Set4")
plt.subplot(3,2,5)
sns.countplot(data=walmart,x='Gender',hue='Marital_Status',order=walmart['Gender'].value_counts().index,palette="Accent")
plt.subplot(3,2,6)
sns.countplot(data=walmart,x='Gender',hue='Product_Category',order=walmart['Gender'].value_counts().index,palette="Paired")
plt.show()

```




- Age group 26-45 have most number of buyers both across males and female
- Male buyers are mostly from occupation 4 and city category B, same hold true for female buyers
- Most of the male and female buyers are unmarried.

In [1075]:

```
sns.countplot(data=walmart,x='Product_Category',hue='Gender',order=walmart['Product_Category'].value_counts().index,palette=
```

Out[1075]:

<Axes: xlabel='Product\_Category', ylabel='count'>



- Most of the females buyers buy product\_category 5 whereas the most of male buyers buy product category 1


In [1076]:

```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                        550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  object
7   Marital_Status                       550068 non-null  object
8   Product_Category                     550068 non-null  int64
9   Purchase                             550068 non-null  int64
dtypes: int64(4), object(6)
memory usage: 42.0+ MB
```

In [1077]:

```
plt.figure(figsize=(20,20))
plt.subplot(3,2,1) ##1 shows the position
sns.countplot(data=walmart,x='Occupation',hue='Age',order=walmart['Occupation'].value_counts().index,palette="Dark2")
plt.subplot(3,2,2)
sns.countplot(data=walmart,x='Age',hue='City_Category',order=walmart['Age'].value_counts().index,palette="Set2")
plt.subplot(3,2,3)
sns.countplot(data=walmart,x='Age',hue='Stay_In_Current_City_Years',order=walmart['Age'].value_counts().index,palette="Set1")
plt.subplot(3,2,4)
sns.countplot(data=walmart,x='Age',hue='Marital_Status',order=walmart['Age'].value_counts().index,palette="Set2")
plt.subplot(3,2,5)
sns.countplot(data=walmart,x='Product_Category',hue='Age',order=walmart['Product_Category'].value_counts().index,palette="Set1")
plt.show()
```




- Buyers in the age group 0-17,18-25 are mostly unmarried and 46-55+ are mostly married
- Product category 5,1,8 are most popular amongst age group 26-25
- Preference of city category for age group 26-35 is B, A,C , whereas for age group 36-45, B,C,A



In [1078]:


```
plt.figure(figsize=(20,20))
plt.subplot(2,2,1) ##1 shows the position
sns.countplot(data=walmart,x='City_Category',hue='Stay_In_Current_City_Years',order=walmart['City_Category'].value_counts())
plt.subplot(2,2,2)
sns.countplot(data=walmart,x='City_Category',hue='Marital_Status',order=walmart['City_Category'].value_counts().index,palette='magma')
plt.subplot(2,2,3)
sns.countplot(data=walmart,x='City_Category',hue='Product_Category',order=walmart['City_Category'].value_counts().index,palette='magma')
plt.show()
```



- City Category B,A has 5 as most bought product category whereas 1 is the most bought product category in city category C

In [1079]:

```
plt.figure(figsize=(20,20))
plt.subplot(2,1,1) ##1 shows the position
sns.countplot(data=walmart,x='Stay_In_Current_City_Years',hue='Marital_Status',order=walmart['Stay_In_Current_City_Years'].value_counts().index)
plt.subplot(2,1,2)
sns.countplot(data=walmart,x='Stay_In_Current_City_Years',hue='Product_Category',order=walmart['Stay_In_Current_City_Years'].value_counts().index)
plt.show()
```



## Bivariate analysis for Categorical-Numerical or Numerical-Categorical

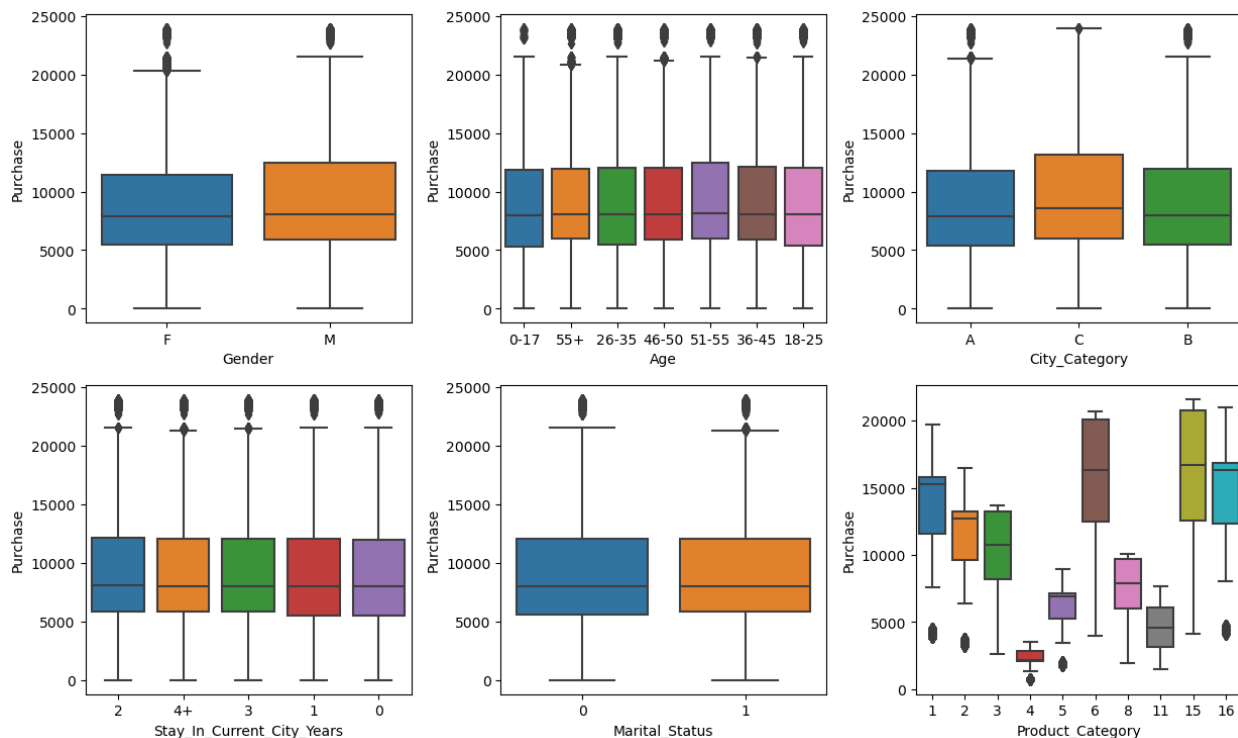
In [1080]:

```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   User_ID                     550068 non-null  int64
1   Product_ID                  550068 non-null  object
2   Gender                       550068 non-null  object
3   Age                          550068 non-null  object
4   Occupation                   550068 non-null  int64
5   City_Category                550068 non-null  object
6   Stay_In_Current_City_Years  550068 non-null  object
7   Marital_Status               550068 non-null  object
8   Product_Category             550068 non-null  int64
9   Purchase                     550068 non-null  int64
dtypes: int64(4), object(6)
memory usage: 42.0+ MB
```

In [1081]:

```
plt.rcParams["figure.figsize"] = [12.50, 7.50]
plt.rcParams["figure.autolayout"] = True
plt.subplot(2,3,1) ##1 shows the position
sns.boxplot(data=walmart,x='Gender',y='Purchase')
plt.subplot(2,3,2)
sns.boxplot(data=walmart,x='Age',y='Purchase')
plt.subplot(2,3,3)
sns.boxplot(data=walmart,x='City_Category',y='Purchase')
plt.subplot(2,3,4)
sns.boxplot(data=walmart,x='Stay_In_Current_City_Years',y='Purchase')
plt.subplot(2,3,5)
sns.boxplot(data=walmart,x='Marital_Status',y='Purchase')
plt.subplot(2,3,6)
sns.boxplot(data=top10_productcategory,x='Product_Category',y='Purchase')
plt.show()
```

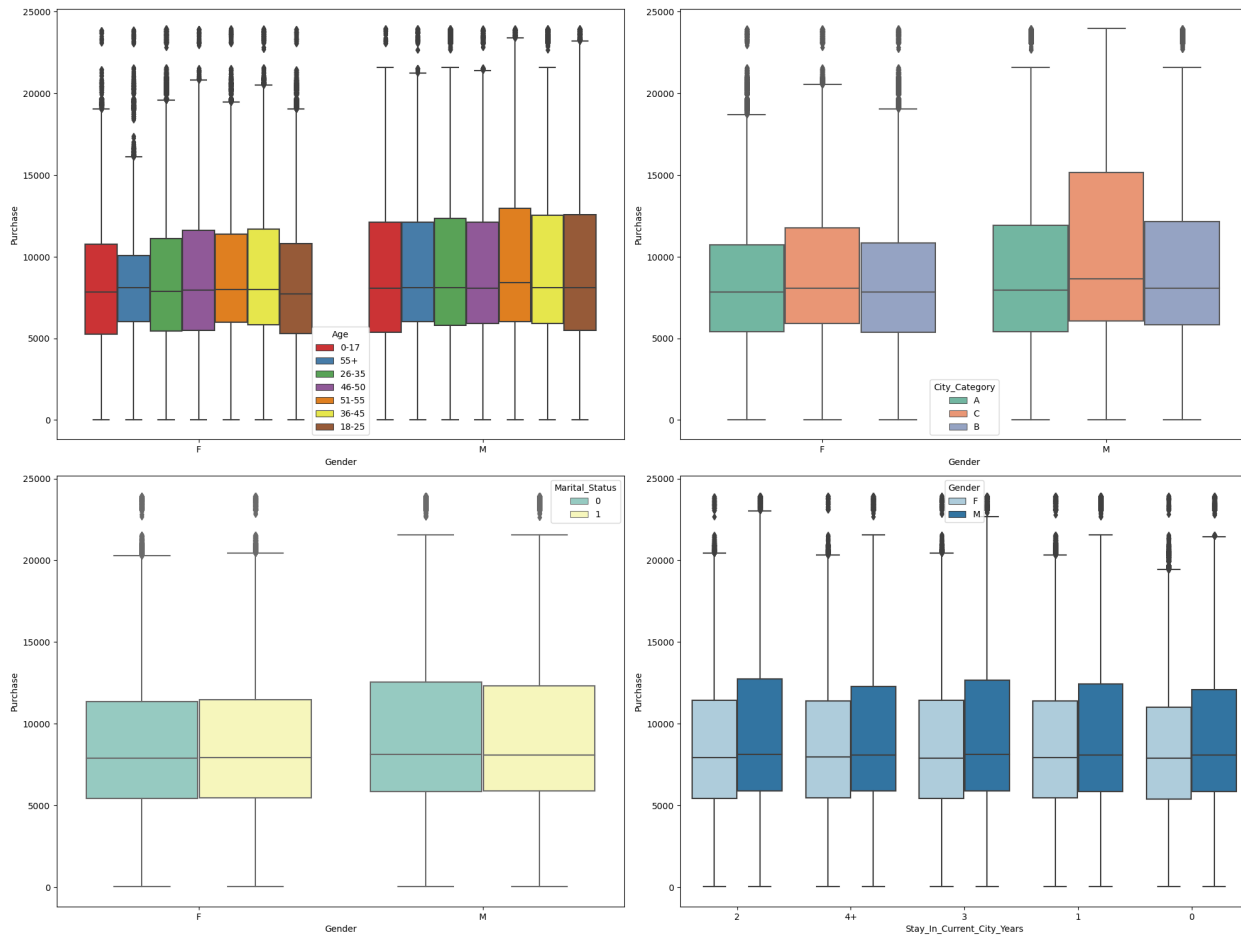


- Median purchase value for male is higher than for female, more number of outliers in female purchases
- Median purchase value of buyer from city category C is higher and has lesser outliers
- Product category 15, 6 have the highest number of buyers and purchase amount is also high with very less outliers

## Multivariate Analysis

In [1082]:

```
fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 15))
plt.rcParams["figure.autolayout"] = True
fig.subplots_adjust(top=1.5)
sns.boxplot(data=walmart, y='Purchase', x='Gender', hue='Age', palette='Set1', ax=axs[0,0])
sns.boxplot(data=walmart, y='Purchase', x='Gender', hue='City_Category', palette='Set2', ax=axs[0,1])
sns.boxplot(data=walmart, y='Purchase', x='Gender', hue='Marital_Status', palette='Set3', ax=axs[1,0])
sns.boxplot(data=walmart, y='Purchase', x='Stay_In_Current_City_Years', hue='Gender', palette='Paired', ax=axs[1,1])
plt.show()
```



- Female purchases has most outliers

## Missing Value & Outlier Detection

In [1083]:

```
walmart.isnull().sum()
```

Out[1083]:

```
User_ID          0
Product_ID       0
Gender           0
Age              0
Occupation       0
City_Category    0
Stay_In_Current_City_Years  0
Marital_Status   0
Product_Category 0
Purchase         0
dtype: int64
```


- No null values in the data

In [1084]:

```
walmart["Purchase"].plot(kind='box')
```

Out[1084]:

&lt;Axes: &gt;



In [1085]:

```
p_25_purchase = np.percentile(walmart['Purchase'],25) ##25th percentile
p_50_purchase = np.percentile(walmart['Purchase'],50) ##50 percentile
p_75_purchase = np.percentile(walmart['Purchase'],75) ##75 percentile
IQR_purchase = p_75_purchase - p_25_purchase
upper_purchase = p_75_purchase + 1.5*IQR_purchase
lower_purchase = max(p_25_purchase - 1.5*IQR_purchase,0)
purchase_outliers=walmart.loc[walmart['Purchase'] > upper_purchase]["Purchase"]
percent_outliers_purchase=(len(purchase_outliers)/len(walmart['Purchase']))*100
print(f"Upper Whisker: {upper_purchase} \nLower Whisker: {lower_purchase}\nIQR: {IQR_purchase}\nOutlier in Purchase:{purchase_outliers}")
```

Upper Whisker: 21400.5

Lower Whisker: 0

IQR: 6231.0

Outlier in Purchase:[23603 23792 23233 ... 23529 23663 23496]

Number of Outliers in Purchase:2677

percent\_outliers\_purchase:0.49

## Answering questions

### Are women spending more money per transaction than men? Why or Why not?

In [1086]:

```
df_purchase=walmart.groupby(['User_ID', 'Gender'])['Purchase'].sum().reset_index()
```

#### Female customers/buyers

In [1087]:

```
df_purchase_female=df_purchase[df_purchase["Gender"]=="F"]
print("Female population mean :",df_purchase_female["Purchase"].mean())
print("Female population standard deviation :",df_purchase_female["Purchase"].std())
```

Female population mean : 712024.3949579832

Female population standard deviation : 807370.7261464577

**Male customers/buyers**

In [1088]:

```
df_purchase_male=df_purchase[df_purchase["Gender"]=="M"]
print("Male population mean :",df_purchase_male["Purchase"].mean())
print("Male population standard deviation :",df_purchase_male["Purchase"].std())
```

Male population mean : 925344.4023668639

Male population standard deviation : 985830.100795388

- Males spend more per transaction than females
- Males earning more than females could be reason behind this behavior , or possibilities are that females prefer buying online

Let us assume the confidence interval to be 95% and desired margin of error to be plus-minus 3

**Sample size calculator**

Confidence Level:

95% ▼

Population Size:

414259

Margin of Error:

3% ▼

Ideal Sample Size:

1065

In [1089]:

```
genders = ["M", "F"]

male_sample_size = 3000
female_sample_size = 1500
num_repetitions = 1000
male_means = []
female_means = []
diff = []


for _ in range(num_repetitions):
    male_mean = df_purchase_male["Purchase"].sample(male_sample_size, replace=True).mean()
    female_mean = df_purchase_female["Purchase"].sample(female_sample_size, replace=True).mean()
    male_means.append(male_mean)
    female_means.append(female_mean)
    diff.append(male_mean-female_mean)
```

In [1090]:

```
fig, axis = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))

axis[0].hist(male_means, bins=35)
axis[1].hist(female_means, bins=35)
axis[0].set_title("Male - Distribution of means, Sample size: 3000")
axis[1].set_title("Female - Distribution of means, Sample size: 1500")

plt.show()
```



In [1091]:

```
print("Population mean - Mean of sample means of amount spend for Male: {:.2f}".format(np.mean(male_means)))
print("Population mean - Mean of sample means of amount spend for Female: {:.2f}".format(np.mean(female_means)))

print("\nMale - Sample mean: {:.2f}, Sample std: {:.2f}".format(df_purchase_male['Purchase'].mean(), df_purchase_male['Purchase'].std()))
print("Female - Sample mean: {:.2f}, Sample std: {:.2f}".format(df_purchase_female['Purchase'].mean(), df_purchase_female['Purchase'].std()))
```

Population mean - Mean of sample means of amount spend for Male: 925731.27  
 Population mean - Mean of sample means of amount spend for Female: 712134.89

Male - Sample mean: 925344.40, Sample std: 985830.10  
 Female - Sample mean: 712024.39, Sample std: 807370.73

### Calculating the 95% confidence interval

In [1092]:

```
z=norm.ppf(.975)
z ## z_score
```

Out[1092]:

1.959963984540054

In [1093]:

```
male_margin_of_error_clt = z*df_purchase_male['Purchase'].std()/np.sqrt(len(df_purchase_male))
male_sample_mean = df_purchase_male['Purchase'].mean()
male_lower_lim = male_sample_mean - male_margin_of_error_clt
male_upper_lim = male_sample_mean + male_margin_of_error_clt

female_margin_of_error_clt = z*df_purchase_female['Purchase'].std()/np.sqrt(len(df_purchase_female))
female_sample_mean = df_purchase_female['Purchase'].mean()
female_lower_lim = female_sample_mean - female_margin_of_error_clt
female_upper_lim = female_sample_mean + female_margin_of_error_clt

diff_margin_of_error_clt = z*(np.std(diff))/np.sqrt(len(diff))
diff_sample_mean = np.mean(diff)
diff_lower_lim = diff_sample_mean - diff_margin_of_error_clt
diff_upper_lim = diff_sample_mean + diff_margin_of_error_clt

print("Male confidence interval of means: ({:.2f}, {:.2f})".format(male_lower_lim, male_upper_lim))
print("Female confidence interval of means: ({:.2f}, {:.2f})".format(female_lower_lim, female_upper_lim))
print("confidence interval of difference of male and female means: ({:.2f}, {:.2f})".format(diff_lower_lim, diff_upper_lim))
```

Male confidence interval of means: (895618.38, 955070.43)  
 Female confidence interval of means: (673255.48, 750793.30)  
 confidence interval of difference of male and female means: (211935.10, 215257.65)

Now we can infer about the population that, **95% of the times**:

1. Average amount spend by male customer will lie in between: **(895617.83, 955070.97)**
2. Average amount spend by female customer will lie in between: **(673254.77, 750794.02)**


**Are confidence intervals of average male and female spending overlapping? How can Walmart leverage this conclusion to make changes or improvements?**

Lets see

In [1094]:

```
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

fig, ax = plt.subplots(figsize=(12, 6))
sns.kdeplot(data=male_means,
            color='crimson', label='male_means', fill=True, ax=ax)
sns.kdeplot(data=female_means,
            color='limegreen', label='female_means', fill=True, ax=ax)
sns.kdeplot(data=diff,
            color='red', label='diff of sample means', fill=True, ax=ax)
ax.legend()
plt.tight_layout()
plt.show()
```



- Neither the means nor the confidence interval are overlapping
- Also as per the article <https://towardsdatascience.com/why-overlapping-confidence-intervals-mean-nothing-about-statistical-significance-48360559900a> (https://towardsdatascience.com/why-overlapping-confidence-intervals-mean-nothing-about-statistical-significance-48360559900a), **If the 95% CI of the difference contains 0, then there is no difference in age between groups. If it doesn't contain 0, then there is a statistically significant difference between groups**
- As it turns out the difference is statistically significant since the 95% CI (shaded red region) doesn't contain 0.
- More lucrative discounts to lure female customers, special promotional offers on International Women's Day , Mother's Day , Daughter's Day, and on their occasions like Anniversay , Birthdays etc



## Are married/unmarried spending more money per transaction than men? Why or Why not?

In [1095]:

```
walmart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
 #   Column                      Non-Null Count  Dtype
---  -
 0   User_ID                    550068 non-null  int64
 1   Product_ID                 550068 non-null  object
 2   Gender                     550068 non-null  object
 3   Age                        550068 non-null  object
 4   Occupation                 550068 non-null  int64
 5   City_Category              550068 non-null  object
 6   Stay_In_Current_City_Years 550068 non-null  object
 7   Marital_Status              550068 non-null  object
 8   Product_Category           550068 non-null  int64
 9   Purchase                   550068 non-null  int64
dtypes: int64(4), object(6)
memory usage: 42.0+ MB
```

In [1096]:

```
df_purchase_maritalstatus=walmart.groupby(['User_ID', 'Marital_Status'])['Purchase'].sum().reset_index()
df_purchase_married=df_purchase_maritalstatus[df_purchase_maritalstatus["Marital_Status"]==1]
print("Married population mean :",df_purchase_married["Purchase"].mean())
print("Married population standard deviation :",df_purchase_married["Purchase"].std())
df_purchase_unmarried=df_purchase_maritalstatus[df_purchase_maritalstatus["Marital_Status"]==0]
print("Unmarried population mean :",df_purchase_unmarried["Purchase"].mean())
print("Unmarried population standard deviation :",df_purchase_unmarried["Purchase"].std())
```

```
Married population mean : 843526.7966855295
Married population standard deviation : 935352.1158252308
Unmarried population mean : 880575.7819724905
Unmarried population standard deviation : 949436.249555238
```

- Unmarried people spend more than married people, unmarried people generally focus less on savings(lesser liabilities) hence they might be spending more

Let us assume the confidence interval to be 95% and desired margin of error to be plus-minus 3

In [1097]:

```
df_purchase_maritalstatus["Marital_Status"].value_counts()
```

Out[1097]:

```
0    3417
1    2474
Name: Marital_Status, dtype: int64
```

In [1098]:

```
unmarried_sample_size = 2000
married_sample_size = 1000
num_repetitions = 1000
unmarried_means = []
married_means = []
diff_ms =[]


for _ in range(num_repetitions):
    unmarried_mean = df_purchase_unmarried["Purchase"].sample(unmarried_sample_size, replace=True).mean()
    married_mean = df_purchase_married["Purchase"].sample(married_sample_size, replace=True).mean()
    unmarried_means.append(unmarried_mean)
    married_means.append(married_mean)
    diff_ms.append(unmarried_mean-married_mean)
```

In [1099]:

```
fig, axis = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))

axis[0].hist(unmarried_means)
axis[1].hist(married_means)
axis[0].set_title("Unmarried - Distribution of means, Sample size: 2000")
axis[1].set_title("Married - Distribution of means, Sample size: 1000")

plt.show()
print("Population mean - Mean of sample means of amount spend for Married: {:.2f}".format(np.mean(married_means)))
print("Population mean - Mean of sample means of amount spend for Unmarried: {:.2f}".format(np.mean(unmarried_means)))
```



Population mean - Mean of sample means of amount spend for Married: 842832.85

Population mean - Mean of sample means of amount spend for Unmarried: 880427.27

### Calculating the 95% confidence interval

In [1100]:

```
unmarried_margin_of_error_clt = z*df_purchase_unmarried['Purchase'].std()/np.sqrt(len(df_purchase_unmarried))
unmarried_sample_mean = df_purchase_unmarried['Purchase'].mean()
unmarried_lower_lim = unmarried_sample_mean - unmarried_margin_of_error_clt
unmarried_upper_lim = unmarried_sample_mean + unmarried_margin_of_error_clt

married_margin_of_error_clt = z*df_purchase_married['Purchase'].std()/np.sqrt(len(df_purchase_married))
married_sample_mean = df_purchase_married['Purchase'].mean()
married_lower_lim = married_sample_mean - married_margin_of_error_clt
married_upper_lim = married_sample_mean + married_margin_of_error_clt

diff_ms_margin_of_error_clt = z*(np.std(diff_ms))/np.sqrt(len(diff_ms))
diff_ms_sample_mean = np.mean(diff_ms)
diff_ms_lower_lim = diff_ms_sample_mean - diff_ms_margin_of_error_clt
diff_ms_upper_lim = diff_ms_sample_mean + diff_ms_margin_of_error_clt

print("Unmarried confidence interval of means: ({:.2f}, {:.2f})".format(unmarried_lower_lim, unmarried_upper_lim))
print("Married confidence interval of means: ({:.2f}, {:.2f})".format(married_lower_lim, married_upper_lim))
print("confidence interval of difference of married and unmarried means: ({:.2f}, {:.2f})".format(diff_ms_lower_lim, diff_ms_upper_lim))
```

Unmarried confidence interval of means: (848741.77, 912409.80)

Married confidence interval of means: (806669.51, 880384.08)

confidence interval of difference of married and unmarried means: (35337.69, 39851.16)

Now we can infer about the population that, **95% of the times**:


1. Average amount spend by unmarried customer will lie in between: **(848741.77, 912409.80)**
2. Average amount spend by married customer will lie in between: **(806669.51, 880384.08)**

**Are confidence intervals of average unmarried and married spending overlapping? How can Walmart leverage this conclusion to make changes or improvements?**

In [1101]:

```
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

fig, ax = plt.subplots(figsize=(12, 6))
sns.kdeplot(data=unmarried_means,
            color='crimson', label='unmarried_means', fill=True, ax=ax)
sns.kdeplot(data=married_means,
            color='limegreen', label='married_means', fill=True, ax=ax)
sns.kdeplot(data=diff_ms,
            color='red', label='diff of sample means of marital status', fill=True, ax=ax)
ax.legend()
plt.tight_layout()
plt.show()
```



- Married and unmarried sample means overlap with each other
- Also as per the article <https://towardsdatascience.com/why-overlapping-confidence-intervals-mean-nothing-about-statistical-significance-48360559900a> (https://towardsdatascience.com/why-overlapping-confidence-intervals-mean-nothing-about-statistical-significance-48360559900a), If the 95% CI of the difference contains 0, then there is no difference in age between groups. If it doesn't contain 0, then there is a statistically significant difference between groups
- As it turns out the difference is **statistically insignificant** since the 95% CI (shaded red region) contain 0.

## Are younger/older spending more money per transaction than men? Why or Why not?

In [1102]:

```
df_purchase_age["Age"].unique()
```

Out[1102]:

```
array(['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25'],
      dtype=object)
```

In [1103]:

```

df_purchase_age=walmart.groupby(['User_ID', 'Age'])['Purchase'].sum().reset_index()
sample_size = 700
num_repetitions = 1000

all_means = {}

age_intervals = ['0-17', '18-25', '26-35', '36-45', '46-50', '51-55', '55+']
for age_interval in age_intervals:
    all_means[age_interval] = []

for age_interval in age_intervals:
    for _ in range(num_repetitions):
        mean = df_purchase_age[df_purchase_age['Age']==age_interval].sample(sample_size, replace=True)['Purchase'].mean()
        all_means[age_interval].append(mean)


```

In [1104]:

```
sns.kdeplot(data=all_means)
```

Out[1104]:

&lt;Axes: ylabel='Density'&gt;



In [1105]:

```

for val in ['26-35', '36-45', '18-25', '46-50', '51-55', '55+', '0-17']:

    new_df = df_purchase_age[df_purchase_age['Age']==val]

    margin_of_error_clt = 1.96*new_df['Purchase'].std()/np.sqrt(len(new_df))
    sample_mean = new_df['Purchase'].mean()
    lower_lim = sample_mean - margin_of_error_clt
    upper_lim = sample_mean + margin_of_error_clt

    print("For age {} --> confidence interval of means: {:.2f}, {:.2f}".format(val, lower_lim, upper_lim))

```

```

For age 26-35 --> confidence interval of means: (945034.42, 1034284.21)
For age 36-45 --> confidence interval of means: (823347.80, 935983.62)
For age 18-25 --> confidence interval of means: (801632.78, 908093.46)
For age 46-50 --> confidence interval of means: (713505.63, 871591.93)
For age 51-55 --> confidence interval of means: (692392.43, 834009.42)
For age 55+ --> confidence interval of means: (476948.26, 602446.23)
For age 0-17 --> confidence interval of means: (527662.46, 710073.17)

```

- Customer between age 18-45 are the potential customer, to retain them , monthly/annual memberships plans can be launched with additional discount and wallet cashback
- Home delivery options to senior citizen (51-55+) so that they buy more

- Survey for Mid age customers can be conducted to understand the mindset and offers they are looking for under product category 5,1,8